

Lecture 1: Introduction to Low-Resource Task Adaption

Dr. LI Yinqiao

li.yin.qiao.2012@hotmail.com

<https://liyinqiao.com>

Presenter Video



JOHN DONNE

(English Poet, 1572-1631)

>> *“No man is an island, entire of itself; every man is a piece of the continent, a part of the main.”*
- *Devotions upon Emergent Occasion (1624)*

*Available in English,
Cantonese and Mandarin*

COURSE PATTERN

- Short Lectures
 - Each lecture is 10-20 minutes long and focuses on one specific topic
 - The lectures cover the main concepts and key knowledge points of the tutorial
- After-Class Exercises
 - Critical reading and reflection on related research
 - Surveying recent advances in related research directions
 - Hands-on coding projects
- Communication
 - GitHub Issues for public discussions, including questions about the tutorials and coding projects hosted under the GitHub Organization.
 - Email is reserved for personal questions

CONTENT

- What Will You Learn from This Tutorial?
 - Understand the basic definition of large language models (LLMs) ★☆☆^[1]
 - Learn the general process of developing and deploying LLMs ★☆☆
 - Recognize the low-resource challenges of applying LLMs to real-world tasks ★☆☆
 - Get a bird's-eye view of low-resource task adaptation ★★☆☆

OUTLINE

- What Are Large Language Models (LLMs)?
 - History and Recent Applications
 - Evolution of LLMs
 - Deployment Pipeline for Downstream Tasks
- Are LLMs Perfect for Every Task?
 - Scarcity of Task-Specific Data
 - Three Key “How” Questions
- Low-Resource Task Adaptation
 - Problem Formulation and Framework
 - Taxonomy of Related Work
- Summary ● Bibliography for Further Study
- Exercises ● Index
- Acknowledgements

[1] The stars indicate the difficulty level of each learning objective.

CONTENT

- What Will You Learn from This Tutorial?
 - Understand the basic definition of large language models (LLMs) ★☆☆^[1]
 - Learn the general process of developing and deploying LLMs ★☆☆
 - Recognize the low-resource challenges of applying LLMs to real-world tasks ★☆☆
 - Get a bird's-eye view of low-resource task adaptation ★★☆☆

OUTLINE

- What Are Large Language Models (LLMs)?
 - History and Recent Applications
 - Evolution of LLMs
 - Deployment Pipeline for Downstream Tasks
- Are LLMs Perfect for Every Task?
 - Scarcity of Task-Specific Data
 - Three Key “How” Questions
- Low-Resource Task Adaptation
 - Problem Formulation and Framework
 - Taxonomy of Related Work
- Summary ● Bibliography for Further Study
- Exercises ● Index
- Acknowledgements

[1] The stars indicate the difficulty level of each learning objective.

CONTENT

- What Will You Learn from This Tutorial?
 - Understand the basic definition of large language models (LLMs) ★☆☆^[1]
 - Learn the general process of developing and deploying LLMs ★☆☆
 - Recognize the low-resource challenges of applying LLMs to real-world tasks ★☆☆
 - Get a bird's-eye view of low-resource task adaptation ★★☆☆

OUTLINE

- What Are Large Language Models (LLMs)?
 - History and Recent Applications
 - Evolution of LLMs
 - Deployment Pipeline for Downstream Tasks
- Are LLMs Perfect for Every Task?
 - Scarcity of Task-Specific Data
 - Three Key “How” Questions
- Low-Resource Task Adaptation
 - Problem Formulation and Framework
 - Taxonomy of Related Work
- Summary ● Bibliography for Further Study
- Exercises ● Index
- Acknowledgements

[1] The stars indicate the difficulty level of each learning objective.

CONTENT

- What Will You Learn from This Tutorial?
 - Understand the basic definition of large language models (LLMs) ★☆☆☆^[1]
 - Learn the general process of developing and deploying LLMs ★☆☆☆
 - Recognize the low-resource challenges of applying LLMs to real-world tasks ★☆☆☆
 - Get a bird's-eye view of low-resource task adaptation ★★☆☆

OUTLINE

- What Are Large Language Models (LLMs)?
 - History and Recent Applications
 - Evolution of LLMs
 - Deployment Pipeline for Downstream Tasks
- Are LLMs Perfect for Every Task?
 - Scarcity of Task-Specific Data
 - Three Key “How” Questions
- Low-Resource Task Adaptation
 - Problem Formulation and Framework
 - Taxonomy of Related Work
- Summary ● Bibliography for Further Study
- Exercises ● Index
- Acknowledgements

[1] The stars indicate the difficulty level of each learning objective.

CONTENT

- What Will You Learn from This Tutorial?
 - Understand the basic definition of large language models (LLMs) ★☆☆^[1]
 - Learn the general process of developing and deploying LLMs ★☆☆
 - Recognize the low-resource challenges of applying LLMs to real-world tasks ★☆☆
 - Get a bird's-eye view of low-resource task adaptation ★★☆☆



OUTLINE

- What Are Large Language Models (LLMs)?
 - History and Recent Applications
 - Evolution of LLMs
 - Deployment Pipeline for Downstream Tasks
- Are LLMs Perfect for Every Task?
 - Scarcity of Task-Specific Data
 - Three Key “How” Questions
- Low-Resource Task Adaptation
 - Problem Formulation and Framework
 - Taxonomy of Related Work
- Summary ● Bibliography for Further Study
- Exercises ● Index
- Acknowledgements

[1] The stars indicate the difficulty level of each learning objective.

What Are Language Models (LMs)?

- Everything starts from language modeling
 - How can a computer understand languages and generate human-like outputs?
 - Traditional computational view: statistical models
 - Human-inspired view: neural models




What Are Language Models (LMs)?

- Everything starts from language modeling
 - How can a computer understand languages and generate human-like outputs?
 - Traditional computational view: statistical models
 - Human-inspired view: neural models
 - With the understanding and generation capabilities, LMs can do more NLP tasks
 - Sentiment analysis
 - Machine translation
 - Information retrieval
 - Even tasks beyond plain text (e.g., speech recognition, handwriting recognition)

Besides Chinese and Western influences, Hong Kong's culture is enriched by its indigenous communities and ethnic minority communities with roots in South and Southeast Asia.


↓ Sentiment Analysis



↓ Machine Translation


除了中國和西方的影響外，香港的文化還因其本土社區和源自南亞和東南亞的少數民族社區而更加豐富多彩。

↓ Information Retrieval




What Are Large Language Models (LLMs)?

- An Explosion of language capability - **large** language models (LLMs)
 - **Large** amounts of data
 - **Large** numbers of neural parameters
 - **Large**-scale computing resources




What Are Large Language Models (LLMs)?

- An Explosion of language capability - **large** language models (LLMs)
 - **Large** amounts of data
 - **Large** numbers of neural parameters
 - **Large**-scale computing resources
- Evolution of LLMs




2020: Breakthrough of LLMs




What Are Large Language Models (LLMs)?

- An Explosion of language capability - **large** language models (LLMs)
 - **Large** amounts of data
 - **Large** numbers of neural parameters
 - **Large**-scale computing resources
- Evolution of LLMs




2020: Breakthrough of LLMs
2020-2023: Scaling up with more data and more parameters




What Are Large Language Models (LLMs)?

- An Explosion of language capability - **large** language models (LLMs)
 - **Large** amounts of data
 - **Large** numbers of neural parameters
 - **Large**-scale computing resources
- Evolution of LLMs




2020: Breakthrough of LLMs
2020-2023: Scaling up with more data and more parameters
2023-present: Shifting toward efficient deployment under computational constraints




What Are Large Language Models (LLMs)?

- An Explosion of language capability - **large** language models (LLMs)
 - **Large** amounts of data
 - **Large** numbers of neural parameters
 - **Large**-scale computing resources
- Evolution of LLMs



2020: Breakthrough of LLMs
2020-2023: Scaling up with more data and more parameters
2023-present: Shifting toward efficient deployment under computational constraints

“Data is not growing, because we have but one Internet. You could even go as far as to say that data is fossil fuel of AI. It was created some how, and now we use it, and we’ve achieved peak data and there’ll be no more^[1]. ”




Ilya Efimovich
Sutskever
(Cofounder and Chief
Scientist of OpenAI)

Projections of the stock of public text and data usage

EPOCH AI [2]

Effective stock (number of tokens)




[1] <https://www.youtube.com/watch?v=1yvBqasHLZs&t>

[2] <https://epoch.ai/publications/will-we-run-out-of-data-limits-of-llm-scaling-based-on-human-generated-data>

What Are Large Language Models (LLMs)?

- An Explosion of language capability - **large** language models (LLMs)
 - **Large** amounts of data
 - **Large** numbers of neural parameters
 - **Large**-scale computing resources
- Evolution of LLMs



2020: Breakthrough of LLMs
2020-2023: Scaling up with more data and more parameters
2023-present: Shifting toward efficient deployment under computational constraints


“Data is not growing, because we have but one Internet. You could even go as far as to say that it is fossil fuel of AI. It will run out some how, and not long. And we’ve achieved peak data and there’ll be no more^[1].”

Sutskever
(Chief Scientist of OpenAI)

So, this is **the end** of the data story?

Projections of the stock of public text and data usage

Effective stock (number of tokens)





EPOCH AI [2]

[1] <https://www.youtube.com/watch?v=1yvBqasHLZs&t>

[2] <https://epoch.ai/publications/will-we-run-out-of-data-limits-of-llm-scaling-based-on-human-generated-data>

What Are Large Language Models (LLMs)?

- Let's take a look at a typical deployment pipeline first
 - Foundation pretraining
 - Instruction tuning
 - Alignment
 - Model adaptation
 - Domain adaptation
 - Task adaptation




Are LLMs Perfect for Every Task? - Not really!

- Let's take a look at a typical deployment pipeline first

- Foundation pretraining
- Instruction tuning
- Alignment
- Model adaptation
 - Domain adaptation
 - Task adaptation
 - In-domain
 - Task-specific
 - Labeled
 - High-quality





Severe data scarcity issue



Are LLMs Perfect for Every Task? - Not really!

- Let's take a look at a typical deployment pipeline first
 - Foundation pretraining
 - Instruction tuning
 - Alignment
 - Model adaptation
 - Domain adaptation
 - Task adaptation
 - In-domain
 - Task-specific
 - Labeled
 - High-quality
- The data scarcity issue becomes **more serious** in rapid explosion of new tasks
 - Emerging programming languages
 - Agent tasks
 - ...



Are LLMs Perfect for Every Task? - Not really!

- Let's take a look at a typical deployment pipeline first

- Foundation pretraining

- The data scarcity issue becomes **more**

- **Low-resource Task Island**


- Scarce task-specific labeled data.
- Limited benchmarks and supervision.
- Underserved needs and domains.

- **Mainland of NLP Progress**

- Abundant data and benchmarks.
- Diverse tasks and general capabilities.
- External knowledge and tools.

No Task Is an Island:
Low-Resource Task Adaptation in LLMs
Connecting low-resource tasks to broader NLP progress through data, training, and inference.

• High-quality



Severe data scarcity issue



Code Translation

ZIG



Are LLMs Perfect for Every Task? - Not really!

- Three key “**how**” questions

- 1 **How** can we obtain and leverage high-quality supervision?
- 2 **How** can we perform effective model updating under limited labeled data?
- 3 **How** can we improve reliability at inference time for low-resource LLMs?

Low-Resource Task Adaptation

- Three key “**how**” questions 1 2 3


- Problem Formulation and Framework

- Data-centric methods

- $\{D_A^{1,t}, \dots, D_A^{n,t}\} = \chi(\{D_S^1, \dots, D_S^n\}; M^{t-1}, \mathbb{D}_R)$

- Number of training tasks: n
- Post-training step: t
- Seed labeled datasets: $\{D_S^1, \dots, D_S^n\}$
- Datasets applied at the t_{th} post-training step: $\{D_A^{1,t}, \dots, D_A^{n,t}\}$
- Model parameters from the last training step: M^{t-1}
- Set of auxiliary unlabeled data resources: $\mathbb{D}_R$
- Data-centric procedure: $\chi(\cdot)$

Large Language Model M^0



- 1 **How** can we obtain and leverage high-quality supervision?


Low-Resource Task Adaptation

- Three key “**how**” questions 1 2 3
- Problem Formulation and Framework

- Data-centric methods
- Effective post-training

- $M^t = \phi(\{D_A^{1,t}, \dots, D_A^{n,t}\}, M^{t-1}; \{M^0, \dots, M^{t-2}\}, \mathbb{M}_R)$

- Number of training tasks: n
- Post-training step: t
- Model parameters from training step t and $t - 1$: M^t and M^{t-1}
- Datasets applied at the t_{th} training step: $\{D_A^{1,t}, \dots, D_A^{n,t}\}$
- Earlier model checkpoints: $\{M^0, \dots, M^{t-2}\}$
- Set of auxiliary models: $\mathbb{M}_R$
- Post-training function: $\phi(\cdot)$



- 2 **How** can we perform effective model updating under limited labeled data?


Low-Resource Task Adaptation

- Three key “**how**” questions ① ② ③
- Problem Formulation and Framework

- Data-centric methods
- Effective post-training
- Reliable inference

$$R = \xi(Q, M^T, r; \mathbb{D}_K)$$


- Number of post-training step: T
- Adapted LLM: M^T
- Query: Q
- Response: R
- Inference strategy: r
- External knowledge or tool base: $\mathbb{D}_K$
- Inference process: $\xi(\cdot)$



③ **How** can we improve reliability at inference time for low-resource LLMs?

Low-Resource Task Adaptation

- Taxonomy of Related Work



Now, Let's Go Back to The Question

So, this is **the end** of the data story?

Definitely not!

- Internet-scale data is being increasingly exhausted for LLM training.
- As a result, leading LLMs are expected to acquire broad world knowledge more comprehensively.
- However, **LLM capabilities remain highly uneven across tasks**. For low-resource tasks, especially newly emerging ones, **achieving effective task adaptation with limited data** still matters for LLM research and development.

Summary

So, this is **the end** of the data story?

Definitely not!

- Internet-scale data is being increasingly exhausted for LLM training.
- As a result, leading LLMs are expected to acquire broad world knowledge more comprehensively.
- However, **LLM capabilities remain highly uneven across tasks**. For low-resource tasks, especially newly emerging ones, **achieving effective task adaptation with limited data** still matters for LLM research and development.

Wrap up of this lecture!

- Understand the basic definition of large language models (LLMs). ★☆☆
- Learn the general process of developing and deploying LLMs. ★☆☆
- Recognize the low-resource challenges of applying LLMs to real-world tasks. ★☆☆
- Get a bird's-eye view of low-resource task adaptation. ★★★



Exercises

Task 1

In what kinds of tasks do LLMs have greater potential than humans? Give some examples. What skills are still difficult for LLMs to master in the near future?

Task 2

Which low-resource tasks are urgent and socially important? Can current LLM systems address them effectively? Survey related open-source projects for one or two such tasks.

Task 3

Based on your background, what machine learning approaches can help alleviate data scarcity in low-resource LLM adaptation?



Please feel free to submit your exercise results and feedback to our GitHub repository (<https://github.com/liyingqiao2012/Low-Resource-Task-Adaptation-in-LLMs/issues>).

exercise

discussion

Bibliography for Further Study

- **LLMs and key stages of its deployment pipeline**
 - An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, and 41 others. 2025a. **Qwen3 technical report**. Preprint, arXiv:2505.09388.
 - Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, and 12 others. 2020. **Language models are few-shot learners**. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc.
 - Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. 2023. **Llama: Open and efficient foundation language models**. Preprint, arXiv:2302.13971.
 - Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh Hajishirzi. 2023d. **Self-instruct: Aligning language models with self-generated instructions**. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 13484–13508, Toronto, Canada. Association for Computational Linguistics.
 - Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2023. **Direct preference optimization: Your language model is secretly a reward model**. In *Advances in Neural Information Processing Systems*, volume 36, pages 53728–53741. Curran Associates, Inc.
 - Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F Christiano, Jan Leike, and Ryan Lowe. 2022. **Training language models to follow instructions with human feedback**. In *Advances in Neural Information Processing Systems*, volume 35, pages 27730–27744. Curran Associates, Inc.
 - Suchin Gururangan, Ana Marasović, Swabha Swayamdipta, Kyle Lo, Iz Beltagy, Doug Downey, and Noah A. Smith. 2020. **Don't stop pretraining: Adapt language models to domains and tasks**. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8342–8360, Online. Association for Computational Linguistics.
 - Jeremy Howard and Sebastian Ruder. 2018. **Universal language model fine-tuning for text classification**. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 328–339, Melbourne, Australia. Association for Computational Linguistics.
 - Vijay Lingam, Atula Neerkaje, Aditya Vavre, Aneesh Shetty, Gautham Krishna Gudur, Joydeep Ghosh, Eunsol Choi, Alex Dimakis, Aleksandar Bojchevski, and Sujay Sanghavi. 2024. **SVFT: parameter-efficient fine-tuning with singular vectors**. In *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*.
 - Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. **Lora: Low-rank adaptation of large language models**. In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net.
 - Xiang Lisa Li and Percy Liang. 2021. **Prefix-tuning: Optimizing continuous prompts for generation**. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing, ACL/IJCNLP 2021, (Volume 1: Long Papers), Virtual Event, August 1-6, 2021*, pages 4582–4597. Association for Computational Linguistics.

Bibliography for Further Study

- **LLMs and key stages of its deployment pipeline**
 - Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego de Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, and 9 others. 2022. **Improving language models by retrieving from trillions of tokens**. In *International Conference on Machine Learning, ICML 2022, 17-23 July 2022, Baltimore, Maryland, USA*, volume 162 of *Proceedings of Machine Learning Research*, pages 2206–2240. PMLR.
 - Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. 2022b. **Chain-of-thought prompting elicits reasoning in large language models**. In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*.
 - Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. **Retrieval-augmented generation for knowledge-intensive nlp tasks**. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474. Curran Associates, Inc.
- **Data Scarcity in task adaptation**
 - Biao Zhang, Zhongtao Liu, Colin Cherry, and Orhan Firat. 2024a. **When scaling meets LLM finetuning: The effect of data, model and finetuning method**. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net.
 - Yizhong Wang, Hamish Ivison, Pradeep Dasigi, Jack Hessel, Tushar Khot, Khyathi Raghavi Chandu, David Wadden, Kelsey MacMillan, Noah A. Smith, Iz Beltagy, and Hannaneh Hajishirzi. 2023c. **How far can camels go? exploring the state of instruction tuning on open resources**. In *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10 - 16, 2023*.
 - Jason Wei, Maarten Bosma, Vincent Y. Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V. Le. 2022a. **Finetuned language models are zero-shot learners**. In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net.
 - Berivan Isik, Natalia Ponomareva, Hussein Hazimeh, Dimitris Paparas, Sergei Vassilvitskii, and Sanmi Koyejo. 2025. **Scaling laws for downstream task performance in machine translation**. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24-28, 2025*. OpenReview.net.
 - Libo Qin, Qiguang Chen, Yuhang Zhou, Zhi Chen, Yinghui Li, Lizi Liao, Min Li, Wanxiang Che, and Philip S. Yu. 2025a. **A survey of multilingual large language models**. *Patterns*, 6(1):101118.
 - Shaolin Zhu, Supryadi, Shaoyang Xu, Haoran Sun, Leiyu Pan, Menglong Cui, Jiangcun Du, Renren Jin, António Branco, and Deyi Xiong. 2024a. **Multilingual large language models: A systematic survey**. *CoRR*, abs/2411.11072.

Index

- [Language Model, 8](#)
- [Language Modeling, 8, 9](#)
- [Statistical Models, 8](#)
- [Neural Models, 8](#)
- [N-Gram LM, 8](#)
- [LMs, 9](#)
- [Sentiment Analysis, 9](#)
- [Machine Translation, 9](#)
- [Information Retrieval, 9](#)
- [Speech Recognition, 9](#)
- [Handwriting Recognition, 9](#)
- [Large Language Models, 10, 11, 12, 13, 14, 15, 26](#)
- [LLMs, 10, 11, 12, 13, 14, 15, 19, 20, 23, 25, 26, 27](#)
- [Word2Vec, 11, 12, 13](#)
- [Distributed Representations, 11, 12, 13](#)
- [Seq2Seq, 11, 12, 13](#)
- [Attention, 11, 12, 13](#)
- [ELMo, 11, 12, 13](#)
- [Pretraining, 11, 12, 13, 16, 17, 18](#)
- [Transformer, 11, 12, 13](#)
- [GPT-1, 11, 12, 13](#)
- [BERT, 11, 12, 13](#)
- [GPT-2, 11, 12, 13](#)
- [T5, 11, 12, 13](#)
- [Scaling Law, 11, 12, 13](#)
- [GPT-3, 11, 12, 13](#)
- [Generative LLMs, 11, 12, 13](#)
- [Instruction Tuning, 12, 13, 16, 17, 18](#)
- [Fine-Tuning, 12, 13](#)
- [InstructGPT, 12, 13](#)
- [Alignment, 12, 13, 16, 17, 18](#)
- [SFT, 12, 13](#)
- [RLHF, 12, 13](#)
- [ChatGPT, 12, 13](#)
- [Chain-of-Thought, 12, 13](#)
- [Chinchilla, 12, 13](#)
- [GPT-4, 12, 13](#)
- [LLaMA, 12, 13](#)
- [Llama 2, 12, 13](#)
- [Foundation Models, 12, 13](#)
- [OpenAI o1, 13](#)
- [Reasoning Model, 13](#)
- [Reinforcement Learning, 13](#)
- [Mixtral, 13](#)
- [MoE, 13](#)
- [DeepSeek-R1, 13](#)
- [Llama 4, 13](#)
- [Multi-Modal LLMs, 13](#)
- [Model Adaptation, 16, 17, 18](#)
- [Domain Adaptation, 16, 17, 18](#)
- [Task Adaptation, 16, 17, 18, 19, 25, 26](#)
- [Data Scarcity, 17, 18, 27](#)
- [Programming Languages, 18](#)
- [Agent, 18](#)
- [Code Translation, 18](#)
- [Low-Resource tasks, 19, 25, 26, 27](#)
- [Labeled Data, 19, 20, 21, 22](#)
- [Low-Resource LLMs, 20, 23](#)
- [Post-Training, 21, 22, 23](#)
- [Unlabeled Data, 21](#)

Acknowledgements

- I would like to **thank** to Prof. SONG Linqi, Prof. XIAO Tong, and other graduate students for their valuable suggestions on the slides and tutorial materials.



Prof. SONG Linqi
City University of
Hong Kong



Prof. XIAO Tong
Northeastern
University




Prof. LIU Jie
North China University
of Technology

Graduate Students:

Ms. YANG Yujie,
Mr. SONG Guangcheng,
Mr. DING Feilong,
Mr. WU Weiquan,
Mr. SHI Jiuan,
Mr. LUO Yingfeng





An Example of Low-Resource Task:
English-Hong Kong Sign Language Machine Translation

[1] The sign-language image is from the Hong Kong Sign Language Association (<https://www.hksla.org.hk/everybody/014.jpg>).